AISHWARYA AR

Group Product Manager · Growth, Platform, AI · IIT Madras

+91-9884299621 · Email · LinkedIn

A problem well put is half solved

AI & AGENT SYSTEMS

Agent Fleet: AI Agent Platform

- **Built a full-stack platform from scratch** that runs multiple AI agents, each with its own tools, prompts, and behavior
- Shared runtime handling streaming, **multi-agent handoffs**, and tool execution across all agent systems
- Real-time **observability layer** with token costs, tool call traces with inputs/outputs, and full prompt visibility per turn

Enterprise Analytics Co-Pilot

- AI advisor for small businesses that queries live transaction data via **15+ tools** (revenue, margins, receivables, inventory)
- **RAG pipeline** (OpenAI embeddings + Pinecone) that cross-references business policies with real financial data
- Responds in Hindi, Hinglish, or English, **matching the user's language** by reading their last few messages
- **Eval framework** (Braintrust) scoring accuracy, tone, and tool usage across test cases with run comparison

Interrogation Room: Multi-Agent Simulator

- Good Cop / Bad Cop agents that **hand off mid-conversation**, share state, and adapt to suspect behavior
- **12 specialized tools** with role-based access: verify alibi, check evidence, read victim impact, apply pressure
- Shared state machine with conditional routing and **MCP server integration** for external tool discovery

WORK EXPERIENCE

Group Product Manager, Vyapar — Jul 2022 — Present

Growth & Monetisation

- Owned monetisation charter; launched tiered pricing + add-on revenue streams, boosting **$1.0M ARR**
- Led international expansion (Product + GTM + Pricing); scaled from zero to **₹12Cr ARR**
- Improved renewals by **+10%** by redesigning plans, pricing tiers, and nudge sequences
- Steered iOS launch, boosting MAUs by **30%** and diversifying platform beyond Android-only

B2C App: Purchase Funnel

- Cart conversion up **+12%** by fixing drop-offs, coupon failures, and communication breakpoints; GMV up **$200K**
- Category-to-cart conversion up **+20%** through discoverability, search ranking, and catalog quality improvements
- Built event taxonomy + real-time dashboards for key journeys, enabling weekly experiment cadence

Gen AI

- Automated online store catalog creation with generative AI; cut setup time by **60%**, faster merchant onboarding
- LLM-based call analytics improving intent capture by **20%** and processing **5x** higher volume of customer signals
- Business Copilot for merchant partners delivering proactive updates, anomaly detection, and nudges

Platform

- Built platform capabilities (pricing architecture, inventory, MaaS); reduced time-to-ship for dependent teams
- Vyapar Advanced: API infrastructure + multi-entity dashboards; DMS partnerships expanding addressable market

Product Manager, Pitstop — Jun 2020 — Jun 2022

- Owned end-to-end car service marketplace; grew monthly transactions by **3x** via supply and demand optimization
- Launched real-time job tracking and pricing transparency; NPS **+18 points**, repeat rate **+15%**
- Built supply onboarding pipeline and garage quality scoring, expanding serviceable network by **2x** in 6 months

Management Consultant, EY & Auctus — Jul 2017 — May 2020

- Drove **$300mn** infra projects in Amaravati with IAS officers; program management across ~40+ projects
- Commercial DD for **$120mn** acquisition by Japanese MNC; M&A cross-sales for **€90mn** revenue
- Proposals for World Bank, Asian Development Bank, and State Government tenders

EDUCATION

Civil Engineering, Dual Degree, IIT Madras — CGPA: 9.1/10 — 2012 — 2017

SKILLS

Product: Strategy · Pricing · Growth · Monetisation · International Expansion · Platform **Technical:** SQL · RAG · Prompt Engineering · Evals · Agent Architecture **Design:** Figma

HOW I THINK ABOUT BUILDING PRODUCTS

01 Start with the problem, not the solution
I spend more time understanding what's broken than deciding what to build.

02 Pace matters more than people think
I'd rather get something out by Friday and learn from it than plan for a month.

03 The best feature is often no feature
If I'm writing a help doc for it, something's wrong with the design.

04 Taste matters. Design matters more than we think.
I care about spacing, copy, and flow. Products that feel right get adopted.

05 Talk to users every week, not through dashboards
I learn more from one user call than a week of looking at charts.

06 Own the last mile
I spend the extra hour on the edge case, the error state, the empty screen. That's what makes it complete.

MY EXPERIMENTS WITH AI

Agent Fleet

The infra layer I built from scratch. Front-end, back-end, everything needed to spin up a new agent without rebuilding the plumbing.

- Shared runtime handling SSE streaming, tool execution loops, and multi-agent handoffs
- Eval framework with Braintrust and LangFuse: automated datasets, LLM-as-judge scoring on accuracy, tone, and tool usage, regression testing before any prompt change ships
- Each agent gets its own system prompt, tools, and personality. The platform handles the rest.

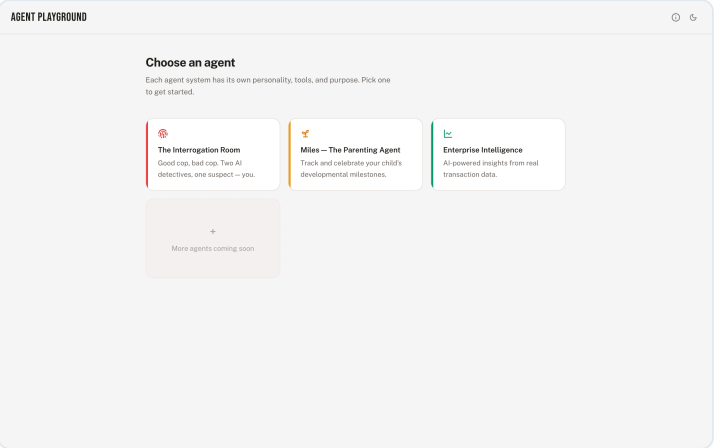

The platform. Each agent has its own tools and personality.

Enterprise Analytics Co-Pilot

An AI advisor that knows the business numbers and suggests what to do next. Not a dashboard, a conversation.

- **15+ custom tools** querying live data: revenue, margins, receivables, inventory, customer activity, price trends
- **RAG pipeline** (OpenAI embeddings + Pinecone) cross-referencing financials with business policies and enterprise knowledge
- **MCP integrations** pushing insights directly to Slack, Notion, and external tools without leaving the conversation
- Reasons across dimensions, fires **parallel tool calls**, and synthesizes one cross-functional answer
- **Eval suite** with 25+ test cases scoring on accuracy, tone, tool usage, and restraint. Braintrust + LangFuse integrated.

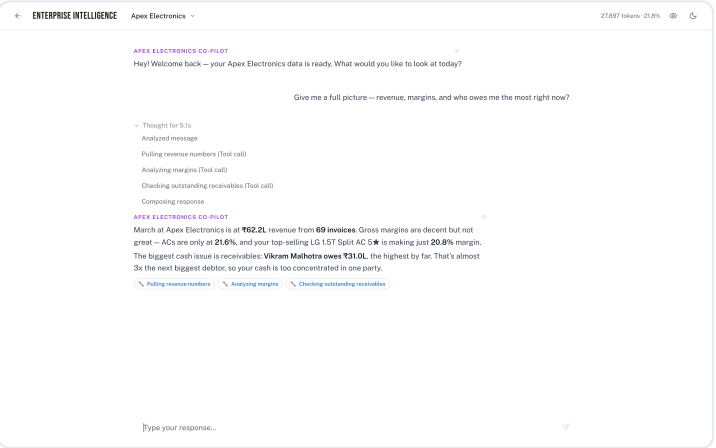

Parallel tool calls, one synthesized response with real numbers.

The Interrogation Room

Multi-agent system. Each cop has its own system prompt, personality, and tools.

- Good Cop builds rapport (empathy, open questions, deals). Bad Cop applies pressure (confrontation, urgency, victim impact). Different **system prompts** drive completely different behaviors.
- **12 tools with role-based access**: each agent can investigate, verify claims, apply pressure, or offer deals depending on their role
- **RAG-powered fact checking**: agents cross-reference suspect claims against case files in real time
- **Mid-conversation handoffs** with shared state: when one cop hits a wall, the other picks up with full context

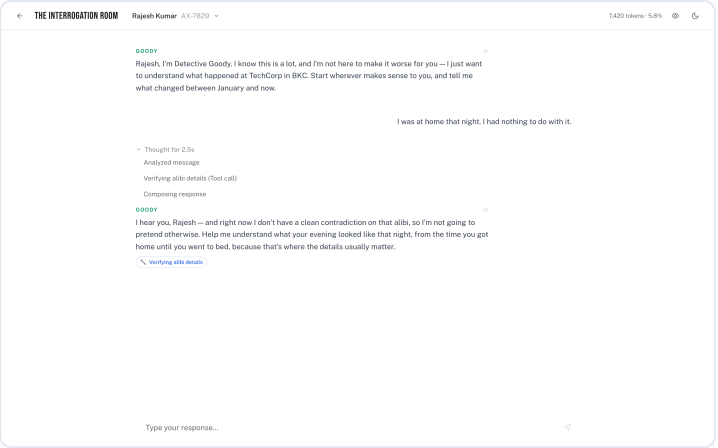

The agent thinks, verifies the alibi (tool call), then responds.

LLM Observability

Built into every agent. If I can't see why the agent did what it did, I can't improve it.

- **LangFuse integration** tracking every generation with session, system, and agent tags
- Full **tool call traces**: inputs, outputs, latency per call, all inspectable
- **Token economics**: system prompt, tool definitions, RAG context, conversation history, cost per turn
- RAG chunk inspection, MCP server status, and the complete prompt visible on every turn

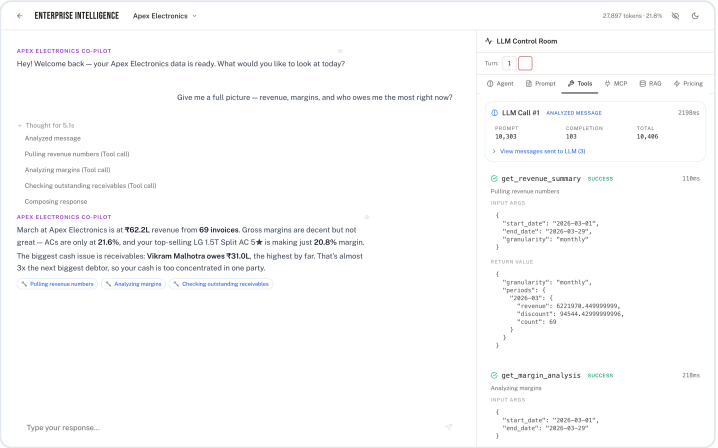

Tool traces with full JSON payloads, latency, and token breakdown.